



FINLY

EXECUTION-REALISTIC CONTROLLED DELEGATION

Positive after next-open execution. Still honest about what is unproven.

Finly returned 15.39% across 415 consumed sessions at next-open fills and five basis points per traded leg. Deterministic code—not an AI model—owns exposure, order fields, maximum loss, and permission to trade.

Bruce Wen · Brandeis University · Alpaca AI Trading Agents Hackathon

Most trading agents collapse judgment, sizing, and execution into one opaque action.

TYPICAL AGENT

A persuasive answer becomes an order.

The same probabilistic component interprets evidence, chooses direction, sizes risk, writes order fields, and explains itself after the fact.

FINLY

Interpretation can inform capital without controlling it.

The model may score bounded evidence, explain uncertainty, or veto. Deterministic code owns exposure, costs, option structure, order fields, and the final permit decision.

Finly's edge is not a larger forecast. It is a smaller, testable trust boundary.

Judges can inspect the evidence instead of taking a demo’s word for it.

CONTROLLED-DELEGATION TRADING AGENT

Finly stayed positive after next-open execution and 25-basis-point cost stress.

The frozen SPY/BIL policy returned +15.39% across 415 consumed sessions under next-open fills and five basis points per traded leg. At a severe 25-basis-point stress, the return remained +10.56%. The model may interpret evidence; deterministic code owns exposure, costs, order fields and the final permission to trade.

Inspect the quantitative evidence

Try the decision record

CONSUMED EXECUTION-REALISM AUDIT

NEXT OPEN · 5 BP / LEG

MODELED TOTAL RETURN	MAXIMUM DRAWDOWN
+15.39% 2025-01-02 — 2026-08-28	-5.45% Annualized volatility 8.12%
The same frozen policy remained positive at +10.56% when the modeled cost rose to 25 basis points per traded leg.	
Adjusted-OHLC theoretical ledger; fractional next-open DAY-order assumption. Consumed retrospective evidence—not a broker fill, options P&L, alpha claim or forecast. SPY returned +33.52% over the same adjusted path.	

VISIBLE ON THE SITE

◆15.39%

NEXT-OPEN · 5 BP PER LEG

415


CONSUMED SESSIONS

◆10.56%

AT 25 BP PER LEG

The same page states that SPY returned 33.52% and that the result is not a broker fill, options P&L, alpha claim, or forecast.

The model can interpret evidence; only deterministic code can authorize capital.



\$366

EXACT MAX LOSS IN CHECKED FIXTURE

4/4


SOURCE-REMOVAL CHALLENGES

32/32

PERTURBATION CHALLENGES

Any mismatch returns NO_TRADE; the model never writes the Alpaca payload.

Returns stayed positive through 25 bp per leg; SPY still won on raw return.



WHAT SURVIVED

◆15.39%

NEXT-OPEN · 5 BP PER LEG

−5.45%

MAXIMUM DRAWDOWN

◆10.56%

AT 25 BP PER LEG

Consumed 2025-01-02 to 2026-08-28 · 415 sessions · adjusted OHLC · fractional next-open DAY-order assumption. SPY returned 33.52% on the same adjusted path.

A \$300 shadow ended at \$351.88 without rounding tiny adjustments into fake fills.

STARTING EQUITY

\$300.00



ENDING EQUITY

\$351.88

+17.29% modeled total return

\$1

MINIMUM ORDER NOTIONAL

\$0.70


SELL-DAY FEE PROXY

12

SUB-MINIMUM ORDERS SKIPPED

Sell-first fractional preview · nine-decimal truncation · one-basis-point traded-leg cost · adjusted-OHLC shadow. Tests affordability and order mechanics; not an Alpaca fill or a promise of future profit.

Finly rejected the backtest that looked best—and kept the weaker-looking production policy.



REJECTED SHADOW · G4

18.97%

annualized vs. SPY 15.11%

Deflated Sharpe probability 3.75% **FAIL**

Worst familywise p-value 0.3718 **FAIL**

Static growth-control independence Unsupported **FAIL**

Authenticated source overlap Not passed **FAIL**

G4 remains a selected-after-history descriptive shadow. It is not the production policy and never received authority to trade.

Attempt 114 turns the next profitability claim into a pre-committed test.

PUBLIC BEFORE THE FIRST SIGNAL

The protocol and executable runtime are already frozen.

A successful public GitHub workflow observed the exact bytes before the exclusive deadline. The rule set forbids skipped sessions, replacement windows, backfill, optional stopping, and a repeat confirmatory test.

GitHub provides a public platform record—not an independent cryptographic timestamp, provider-origin proof, broker fill, or profitability result.

17/17

RUNTIME-BOUND FILES MATCHED

23

FIXED PUBLIC GET CHECKS

254

TIMELY SIGNAL ANCHORS REQUIRED

252

RECONCILED SETTLEMENTS
REQUIRED

TODAY: 0 anchors · 0 settlements
Performance inference and broker mutation remain disabled.

Finly gives judges evidence they can verify—and a future claim they can falsify.

P&L EVIDENCE

+15.39% at next-open fills and 5 bp per leg; +10.56% at 25 bp. SPY returned +33.52%.

TECHNOLOGY

An implemented authority boundary keeps models away from exposure, order fields, maximum loss, and broker permission.

ORIGINALITY

Finly makes correction, rejection, and NO_TRADE observable product behavior—not a disclaimer added after the trade.

FORWARD PROOF

Attempt 114 freezes 17 runtime files and requires 254 anchors plus 252 reconciled settlements before inference.

The proposition is simple: let AI interpret more, authorize less, and earn every stronger profitability claim in public.