

Finly: The AI Trading Agent That Shows Its Work

Bruce Wen · Brandeis University · bwen412@brandeis.edu

Proposal for the Alpaca AI Trading Agents Hackathon · 31 August 2026

Finly is an AI trading agent that shows its work before it risks a dollar. Most trading agents ask one model to read the market, choose a position, size it, and send the order. Finly separates those jobs. AI reviews public information and explains what supports or weakens a trade; fixed, tested code controls the direction, contracts, position size, maximum loss, and every field sent to Alpaca. The AI may lower confidence or stop a trade, but it cannot increase the amount at risk. Investors get the speed of AI without handing a language model unrestricted control of a brokerage account.

For this hackathon, Finly runs in the cloud against a dedicated, verified \$100,000 Alpaca paper account, which follows real prices with virtual money. The product has two parts. Finly Core is the base portfolio: 48.5 percent in QQQ, a fund that tracks the Nasdaq-100; 48.5 percent divided across three sector funds chosen by a transparent price-trend rule; and 3 percent in cash. Finly Options is a separate SPY assistant. SPY tracks the S&P 500. Its general code can build rising- and falling-price two-option plans with fixed maximum loss, while the live competition rules allow only an approved rising-price trade or no trade. We locked the portfolio and rules before the test began, so the account could not be revised after seeing the result [1].

Every live options decision follows the same visible path. Finly gathers current market prices, option quotes, and public Alpaca news. A hosted Qwen3-32B model reads the news and identifies support, conflict, and uncertainty. Code then chooses the direction and two-option trade, fixes the quantity and maximum loss, and runs the required checks: data must be current, the options must have enough market activity to trade, the payoff must add up, the account must have enough buying power, and the decision must remain stable when a source is removed or an important input changes. If any check fails, Finly stops. If they all pass, the approval is tied to one exact order, sent through Alpaca's official connection, and checked against the paper account afterward. The public dashboard then shows a safe account summary and a saved decision record explaining why Finly traded—or why it did not [2,3].

The numbers give judges a reason to test that architecture. In a historical simulation from January 2013 through August 2026, with modeled trading costs, \$10,000 became \$106,711 with Finly Core and \$68,082 with SPY—a \$38,629 difference in ending wealth. A simpler industry version was checked on 21,218 earlier market days. It averaged 13.37 percent growth per year against 9.48 percent for the market and stayed ahead in all twenty-one tests that changed which trading day of the month the holdings were updated. The interactive options demonstration made the risk concrete: one illustrative SPY plan had a \$366 maximum loss and \$634 maximum gain, and reached the same decision after all four sources were removed one at a time and after thirty-two input changes. No broker order or fill occurred. The public suite ran 809 tests: 807 passed, none failed, and two were skipped [2,4]. At the first 4:00 p.m. close, Finly Core was up \$95.32 while SPY was down \$57.99 from the same starting point—a \$153.31 advantage after 15 broker fill events and with no deposits or withdrawals. No options position was open at that close [2].

Finly's advantage is not asking an AI to guess harder. It is making every trade earn approval. That gives investors, educators, and brokerage teams clear answers: what did the AI believe, how much could the trade lose, and what did Alpaca accept? It also meets the hackathon on its own terms: meaningful AI logic, explicit risk checks, the required options component, official Alpaca infrastructure, a dedicated competition account, and visible paper-account results. Judges can open the working web app, replay both the trade and no-trade demonstrations, watch the cloud account, inspect the public GitHub repository, and follow the same story in the demo video and pitch deck [5]. Finly does not merely produce a market opinion; it turns that opinion into a checked, limited, and inspectable action.

References

[1] Finly, [Frozen competition protocol](#) and [deployment record](#). [2] Finly, [Quantitative evidence](#), [earlier-era replay](#), [options decision receipt](#), and [first-close measurement](#). [3] Alpaca, [official MCP server](#); Finly, [cloud runner design](#). [4] Finly, [public automated-test run](#). [5] Finly, [live application](#), [public repository](#), [demo video](#), and [pitch deck](#).