

Finly: The AI Trading Agent That Shows Its Work

Bruce Wen · Brandeis University · bwen412@brandeis.edu
Proposal for the Alpaca AI Trading Agents Hackathon · 31 August 2026

Finly is an AI trading agent that shows its work before it risks a dollar. Most trading agents ask one model to read the market, choose a position, size it, and send the order. Finly separates those jobs. AI interprets public evidence and explains support or conflict; code controls direction, contracts, position size, maximum loss, and every field sent to Alpaca. Investors get the speed of AI without handing a language model unrestricted control of a brokerage account.

For this hackathon, Finly runs in the cloud against a dedicated, verified \$100,000 Alpaca paper account. A frozen four-fund equity allocation provides the portfolio base: QQQ plus three sector funds selected before the opening bell by a transparent momentum rule, with three percent left in cash. The competition's required options component is a separate SPY agent. Its general compiler can construct bullish and bearish debit spreads, while the live contest guard deliberately authorizes only a bullish spread or no trade. Because the allocation and rules were frozen before the scoring window, the forward account records what the system actually does rather than a strategy revised after seeing the result [1].

Every options decision follows the same visible path. Finly first gathers time-stamped market prices, option quotes, scheduled events, and public Alpaca news. A hosted Qwen3-32B model reads the public news and identifies support, conflict, and uncertainty. Code then combines the allowed evidence, chooses the direction and spread, fixes the quantity and maximum loss, and runs the risk gates required before trading: data must be fresh, the options must be liquid, the payoff must add up, the account must have enough buying power, and the decision must remain stable when a source is removed or an important input changes. If any check fails, Finly stops. If they all pass, the approval is tied to one exact order, sent to the Alpaca paper-trading API through Alpaca's official MCP server, and read back from Alpaca. The public dashboard then shows the sanitized account result and a receipt explaining why Finly traded—or why it did not [2,3].

The numbers give judges a reason to test that architecture. In a cost-adjusted historical simulation from January 2013 through August 2026, \$10,000 became \$106,711 with Finly's G4 rule and \$68,082 with SPY—a \$38,629 difference in ending wealth. A fixed replay on 21,218 earlier market days produced 13.37 percent annualized growth against 9.48 percent for the market and stayed ahead across all twenty-one possible offsets of its rebalance schedule. The options demonstration showed the risk in concrete dollars: one SPY debit spread had a \$366 maximum loss and \$634 maximum gain, and survived all four source-removal checks and all thirty-two input changes. The public suite ran 809 tests: 807 passed, none failed, and two were skipped [2,4]. At the first close, Finly was up \$95.32 while same-clock SPY was down \$57.99—a \$153.31 advantage after 15 fills and with no external cashflows [2].

Finly's advantage is not asking an AI to guess harder. It is making every trade earn permission. That gives investors, educators, and brokerage teams clear answers: what did the AI believe, how much could the trade lose, and what did Alpaca accept? It also meets the hackathon on its own terms: meaningful AI logic, explicit risk gates, mandatory options trading, official Alpaca infrastructure, a dedicated competition account, and visible paper P&L. Judges can open the working web app, change the evidence to trigger both trade and no-trade paths, watch the cloud account, inspect the public GitHub repository, and follow the same story in the demo video and pitch deck [5]. Finly does not merely produce a market opinion; it turns that opinion into a checked, limited, and inspectable action.

References

[1] Finly, [Frozen competition protocol](#) and [deployment record](#). [2] Finly, [Quantitative evidence](#), [earlier-era replay](#), [options decision receipt](#), and [first-close measurement](#). [3] Alpaca, [official MCP server](#); Finly, [cloud runner design](#). [4] Finly, [public automated-test run](#). [5] Finly, [live application](#), [public repository](#), [demo video](#), and [pitch deck](#).