



Finly: When a Good Backtest Should Still Not Trade

Bruce Wen · Brandeis University · bwen412@brandeis.edu

Most trading-agent designs collapse two different questions into one. The first is epistemic: can a system form a plausible view about what markets may do? The second is institutional: has that view earned authority to risk capital? Finly separates them. AI may score bounded evidence, provide a rationale, and reduce or veto a deterministic proposal. A code-owned challenge suite may remove evidence families and rerun the checks. The model does not determine direction, horizon, option structure, maximum loss, or order fields. Deterministic code owns those decisions, together with equality checks and the final PERMIT or NO_TRADE outcome. The central claim is therefore narrower than a forecasting claim: useful judgment need not imply capital-bearing authority.

The strongest objection is that this division may remove the very discretion that makes an agent valuable. Momentum effects have substantial precedent [1], and volatility management shows that exposure can be adjusted systematically in response to measured risk [2]. A capable model might appear able to combine such signals with contextual evidence more flexibly than fixed rules. If the model is already being consulted, restricting it from selecting structure or exposure can look unnecessarily conservative.

That objection has force, but it moves too quickly from informational value to authorization. A model can produce an economically coherent rationale while remaining unstable, overfit, or insufficiently tested against alternatives. Historical attractiveness is not the same as permission to deploy. Finly's research record illustrates the distinction. One descriptive candidate allocated 50% to QQQ and divided the remaining 50% equally among the top three sector ETFs ranked by the log return from 252 to 126 sessions before each 21-session rebalance. Replayed from 2 January 2013 through 27 August 2026 on adjusted closes, it returned 967.11% in total, or 18.97% annualized, versus SPY at 580.82% and 15.11%. Modeled volatility was 18.01% versus 16.79%, maximum drawdown was -28.99% versus -33.72%, and annualized notional turnover was 3.78x versus 0.22x. With a five-basis-point one-way cost applied to traded notional, the sign of its historical G4-minus-SPY edge remained positive across all 21 rebalance-anchor offsets and at the tested 5, 10, and 25 basis-point cost levels.

Those figures are attractive, but they did not justify promotion. The candidate's Deflated Sharpe probability was 3.75% against a 95% threshold [3]; its worst adjusted familywise p-value was 0.3718 against 0.05; the static growth-control independence gate was unsupported; and the authenticated source-overlap gate did not pass. Seven later hash-frozen but fully retrospective G6 challengers were assessed, and none was selected on either the primary SPY or growth-control track. Finly conservatively counts 113 effective research attempts, although that count includes controls, rejected or unexecuted suggestions, invalidated runs, an aborted attempt, and reruns. It is not evidence of 113 independent viable strategies, but it does reinforce the need to treat apparently strong results cautiously [4].

G4 is also not the policy Finly would trade. The frozen production book is a lower-risk SPY/BIL trend-and-volatility rule. In its now-consumed fixed 2025–2026 holdout, that policy returned 11.13% annualized against SPY's 19.19%, while reducing annualized volatility from 17.33% to 8.31% and maximum drawdown from -18.76% to -5.79%. It exchanged return for protection; it did not establish that Finly is more likely than not to beat SPY next month.

The methodology also should not be described as fully preregistered. The repository records a local hash freeze of the core formula, partitions, and costs before the first G4 output, but the local timestamp is not an independent time authority. The excess-Sharpe selection rule and later inference corrections changed after the initial output, and every market interval used in the reported analysis is now consumed. That weakens any claim that the exercise constitutes a clean prospective test. The appropriate response is not to hide the limitation, but to narrow the inference drawn from it.

The options component should be interpreted similarly. The historical ETF allocation is not an options profit-and-loss record. Finly's options contribution is the controlled compiler and defined-risk authorization boundary: deterministic code owns direction, horizon, option structure, maximum loss, order fields, equality checks, and the final decision, while AI remains confined to bounded evidentiary judgment. The locked evidence surface records a successful authenticated read-only Alpaca paper-account check; no order or fill is presented as performance evidence [5,6].

Forward Trial 1 defines the next proof stage. A clean clone can verify its local zero-row state, but the record begins with no signal commitments and no settled outcomes. The protocol separates what must be fixed before a decision from what can be measured only afterward. Because a local timestamp and hash cannot prove when external data were observed, production writes, broker authority, and performance inference remain disabled until an independent pre-execution anchor exists. Finly should therefore be judged on the narrower claim it can support today: whether this division of authority is technically coherent, auditable, and actually enforced. Profit remains an empirical question for the forward record.

References. [1] Moskowitz, Ooi, and Pedersen, "Time Series Momentum," *Journal of Financial Economics* 104(2), 2012, [doi:10.1016/j.jfineco.2011.11.003](https://doi.org/10.1016/j.jfineco.2011.11.003). [2] Moreira and Muir, "Volatility-Managed Portfolios," *Journal of Finance* 72(4), 2017, [doi:10.1111/jofi.12513](https://doi.org/10.1111/jofi.12513). [3] Bailey and López de Prado, "The Deflated Sharpe Ratio," *Journal of Portfolio Management* 40(5), 2014, [doi:10.3905/jpm.2014.40.5.094](https://doi.org/10.3905/jpm.2014.40.5.094). [4] White, "A Reality Check for Data Snooping," *Econometrica* 68(5), 2000, [doi:10.1111/1468-0262.00152](https://doi.org/10.1111/1468-0262.00152). [5] Alpaca, "Options Trading," official documentation, docs.alpaca.markets. [6] Finly, "Submission Quantitative Evidence Surface," 2026, claim registry: public/data/submission_claims_lock.json.