

# Finly: A Trading Agent That Knows When Not to Trade

Bruce Wen · Brandeis University · bwen412@brandeis.edu

Proposal for the Alpaca AI Trading Agents Hackathon · Evidence through 30 August 2026

---

What should an AI trading agent be allowed to decide? The obvious answer is that a capable model should read more evidence, form a better view, and place the corresponding trade. Finly begins from the opposite concern. In markets, a persuasive explanation and an authorized exposure are different things: one can be evaluated in words; the other can lose capital before anyone has resolved whether the explanation was sound. Finly therefore gives AI a serious but limited role. A model may interpret a bounded evidence bundle, state its reasons, and veto a proposed action. Deterministic code retains the decisions that directly create risk: market direction, horizon, structure, order fields, maximum loss, and the final choice between permission and no trade. The proposal is not to make the model decorative. It is to make its authority proportionate to the evidence. That division makes the demo inspectable. A judge can follow a source from ingestion to interpretation, compare the model's typed response with the code-owned intent, and see why disagreement removes permission rather than being smoothed into a more confident order. The resulting record identifies which judgment failed and who still owned the risk.

The project tests this boundary against its own strongest temptation. In the consumed, post-selected 2013-01-02–2026-08-27 retrospective replay with modeled 5 bp one-way costs, G4 returned +967.11% versus SPY +580.82%; promotion was rejected because the Deflated Sharpe probability was 3.75% and the worst familywise-adjusted p-value was 37.18%. That is the result a conventional pitch would enlarge and leave unexplained. Finly instead keeps the rejection beside the return. G4 is not the production policy and not evidence of future market superiority [3–5]. The refusal matters because a backtest is most dangerous when it is attractive enough to silence the question of how it was found.

The frozen production policy makes a smaller claim. Production v1 is the frozen unlevered SPY/BIL policy targeting 10% annualized volatility: in the consumed 2025-01-02–2026-08-28 modeled next-open study it returned +15.39% at 5 bp per traded leg and +10.56% at 25 bp, versus SPY +33.52%; at 5 bp its modeled annualized volatility was 8.12% and maximum drawdown was -5.45%, so it was risk-controlled but not market-beating on total return. The study is a theoretical adjusted-OHLC ledger, not broker fills or options profit and loss [1,2,5]. This distinction turns an apparent weakness into useful information. Judges can see what the policy bought, what it surrendered, and what remains unproved.

Finly's final step is prospective rather than rhetorical. Attempts 115 and 116 are publicly registered future-only tests. As of 2026-08-30T08:10:52.000Z, each had zero observed outcomes, and neither supports a performance claim. [6,7] That empty record is intentional: it prevents a favorable history from being relabeled as a forward test. At the Generation 7 capture, 20 projects were visible and zero supplied an exact same-panel submitted-options comparator; missing P&L is unknown, never zero, and supports neither a return matchup nor a competitor rank. [5] Finly should therefore be judged on what it presently demonstrates: a working separation between interpretation and authorization, a research process capable of rejecting its best-looking historical survivor, a frozen policy described without disguising its tradeoff, and a public experiment that can earn a stronger claim later. The governing principle is modest but consequential: when capital is at risk, intelligence should widen inquiry before it widens authority.

## References

- [1] T. J. Moskowitz, Y. H. Ooi, and L. H. Pedersen, "Time Series Momentum," *Journal of Financial Economics* 104 (2012), 228–250, [doi:10.1016/j.jfineco.2011.11.003](https://doi.org/10.1016/j.jfineco.2011.11.003). [2] A. Moreira and T. Muir, "Volatility-Managed Portfolios," *Journal of Finance* 72 (2017), 1611–1644, [doi:10.1111/jofi.12513](https://doi.org/10.1111/jofi.12513). [3] H. White, "A Reality Check for Data Snooping," *Econometrica* 68 (2000), 1097–1126, [doi:10.1111/1468-0262.00152](https://doi.org/10.1111/1468-0262.00152). [4] D. H. Bailey and M. López de Prado, "The Deflated Sharpe Ratio," *Journal of Portfolio Management* 40 (2014), 94–107, [doi:10.3905/jpm.2014.40.5.094](https://doi.org/10.3905/jpm.2014.40.5.094). [5] Finly, [Quantitative Release Gate](#). [6] Finly, [Attempt 115 registration](#). [7] Finly, [Attempt 116 registration](#).