Finly: AI Explains the Trade. Code Controls the Money.

Bruce Wen · Brandeis University · bwen412@brandeis.edu

Proposal for the Alpaca AI Trading Agents Hackathon · 31 August 2026

Finly is a paper-trading bot that lets AI study the market without giving it control of the account. Paper trading follows real prices with virtual money. In Finly's first session on a dedicated $100,000 Alpaca paper account, the portfolio gained $95.32 while SPY—the fund commonly used to represent the S&P 500—lost $57.99 from the same starting point. Finly therefore finished the session $153.31 ahead at the same timestamp. Alpaca recorded fifteen ETF fill events as the allocation was assembled, with no deposits, withdrawals, or options held at that close [1]. One session cannot establish future performance, but it does show that Finly moved from an idea to a running system with a result that judges can inspect.

Finly is one autonomous competition strategy with two coordinated sleeves. The four-fund sleeve keeps 48.5 percent in QQQ, which tracks the Nasdaq-100; divides another 48.5 percent among three sector funds with stronger longer-term price trends; and holds 3 percent in cash. The coordinated options sleeve evaluates one small, capped-loss trade tied to SPY. Fresh price momentum decides whether the short-horizon idea is positive or negative; ordinary options-market evidence and Qwen3-32B's reading of public Alpaca news may support that view, weaken it, or stop it. The model cannot choose the contracts, size, loss limit, or broker fields. Fixed code does that work and limits a live entry to exactly one contract and no more than $500 of possible loss, with a tighter 0.5 percent account-risk budget when applicable [2].

Every decision by the options sleeve follows the same visible path. Finly records current prices, option quotes, trading activity, and public Alpaca news. Code first asks whether the longer-horizon account remains invested enough to justify a short-horizon options hedge or opportunity. It then builds a bullish call spread, a bearish put spread, or no trade; checks price freshness, liquidity, payoff arithmetic, account capacity, and modeled value; and records how the idea behaves when evidence is removed or inputs are nudged. Those challenges diagnose confidence, while the hard safety checks still decide whether money may move. An approved trade is bound to one exact order, saved before mutation, and reconciled with Alpaca afterward, so a restart cannot silently duplicate it [1,3].

The historical evidence provides a longer test of the underlying four-fund research rule, before the competition's 3 percent cash scaling. In a January 2013–August 2026 simulation after modeled trading costs, $10,000 became $106,711 with Finly and $68,082 with SPY, a $38,629 difference in ending wealth. A simpler version built from long-running industry data was tested on 21,218 market days from 1927 through 2007. It averaged 13.37 percent growth per year versus 9.48 percent for the market, remained ahead across all twenty-one tested monthly rebalance schedules, and stayed ahead when the assumed trading cost increased fivefold. The options path was separately recalibrated after a published no-trade diagnostic. Across 517 sampled SPY signal windows, 11 cleared every alpha gate on a symmetric modeled quote surface: seven bullish call spreads and four bearish put spreads. Their maximum losses were $440–$455, conservative modeled values after costs were $10.08–$22.26, and reward-to-risk ratios were 2.30–2.41. That is a bidirectional eligibility test, not options P&L. The full verification run found 826 automated tests: 824 passed, none failed, and two optional private-ledger checks were skipped [2,4].

Finly does not try to win by asking an AI to sound more certain. Its coordinated sleeves make every capital-bearing decision pass through software that can explain the evidence, limit the loss, and leave the account untouched. The four-fund allocation did not change after its paper result, and the revised options policy does not rewrite earlier decisions or P&L: it has its own dated record, code hash, tests, and future broker receipts [2]. An ordinary visitor can see what Finly bought and whether it made money; a technical judge can trace the allocation through its fills and separately inspect a bullish, bearish, or no-trade options decision. That combination—visible performance, two-sided options logic, meaningful AI analysis, and code-owned execution authority—is Finly's case for the Options Alpha Agents track [5].

References

[1] Finly, first-close measurement, competition deployment record, and latest options decision receipt. [2] Finly, frozen allocation protocol, current dated options-policy revision, and quantitative evidence. [3] Alpaca, official MCP server; Finly, cloud runner design. [4] Finly, public automated-test run. [5] Finly, live application, public repository, mathematical technical note, engineering appendix, demo video, and pitch deck.