

Finly: AI Explains the Trade. Code Controls the Money.

Bruce Wen · Brandeis University · bwen412@brandeis.edu

Proposal for the Alpaca AI Trading Agents Hackathon · 2 September 2026

Finly begins with a measurable result. In a January 2013–August 2026 historical simulation after modeled trading costs, the same \$10,000 became \$106,711 with Finly's underlying four-fund rule and \$68,082 with SPY, the fund commonly used to represent the S&P 500. That is a \$38,629 difference in ending wealth. The comparison is historical rather than live, but it answers the first question a judge should ask: whether the portfolio idea produced a result worth turning into a working product.

The answer was then tested in deployment. Finly is a paper-trading bot that lets AI study the market without giving it control of the account. Paper trading follows real prices with virtual money. Through the September 2 close, Finly gained \$141.24 in its dedicated \$100,000 Alpaca paper account while SPY lost \$284.76 from the same starting point. Finly therefore finished \$426.00 ahead at the same 4:00 p.m. timestamp. Alpaca recorded fifteen ETF fill events as the allocation was assembled, with no deposits, withdrawals, or options held at that close [1]. The live window is short, but it shows that Finly moved from a historical research result to a running system with a score judges can inspect.

Finly is one autonomous competition strategy with two coordinated sleeves. The four-fund sleeve keeps 48.5 percent in QQQ, which tracks the Nasdaq-100; divides another 48.5 percent among three sector funds with stronger longer-term price trends; and holds 3 percent in cash. The coordinated options sleeve evaluates one small, capped-loss trade tied to SPY. Fresh price momentum decides whether the short-horizon idea is positive or negative; ordinary options-market evidence and Qwen3-32B's reading of public Alpaca news may support that view, weaken it, or stop it. The model cannot choose the contracts, size, loss limit, or broker fields. Fixed code does that work and limits a live entry to exactly one contract and no more than \$500 of possible loss, with a tighter 0.5 percent account-risk budget when applicable [2].

Every decision by the options sleeve follows the same visible path. Finly records current prices, option quotes, trading activity, and public Alpaca news. Code first asks whether the longer-horizon account remains invested enough to justify a short-horizon options hedge or opportunity. It then builds a bullish call spread, a bearish put spread, or no trade; checks price freshness, liquidity, payoff arithmetic, account capacity, and modeled value; and records how the idea behaves when evidence is removed or inputs are nudged. On September 2 it completed twenty-four live evaluation cycles: fourteen candidates failed certification, six lacked enough model evidence, and four arrived after the entry window. Every cycle ended no trade, so the options sleeve added no new risk [1]. An approved trade would still be bound to one exact order, saved before mutation, and reconciled with Alpaca afterward, preventing a restart from silently duplicating it [3].

The historical evidence goes beyond that headline window. A simpler version built from long-running industry data was tested on 21,218 market days from 1927 through 2007. It averaged 13.37 percent growth per year versus 9.48 percent for the market, remained ahead across all twenty-one tested monthly rebalance schedules, and stayed ahead when the assumed trading cost increased fivefold. The options path was separately recalibrated after a published no-trade diagnostic. Across 517 sampled SPY signal windows, 11 cleared every alpha gate on a symmetric modeled quote surface: seven bullish call spreads and four bearish put spreads. Their maximum losses were \$440–\$455, conservative modeled values after costs were \$10.08–\$22.26, and reward-to-risk ratios were 2.30–2.41. That is a bidirectional eligibility test, not options P&L. The full verification run found 827 automated tests: 825 passed, none failed, and two optional private-ledger checks were skipped [2,4].

Finly does not try to win by asking an AI to sound more certain. Its coordinated sleeves make every capital-bearing decision pass through software that can explain the evidence, limit the loss, and leave the account untouched. The four-fund allocation did not change after its paper result, and the revised options policy does not rewrite earlier decisions or P&L: it has its own dated record, code hash, tests, and future broker receipts [2]. An ordinary visitor can see what Finly bought and whether it made money; a technical judge can trace the allocation through its fills and separately inspect a bullish, bearish, or no-trade options decision. That combination—visible performance, two-sided options logic, meaningful AI analysis, and code-owned execution authority—is Finly's case for the Options Alpha Agents track [5].

References

- [1] Finly, [September 2 same-clock measurement](#), [live options decision record](#), and [competition deployment record](#). [2] Finly, [frozen allocation protocol](#), [current dated options-policy revision](#), and [quantitative evidence](#). [3] Alpaca, [official MCP server](#); Finly, [cloud runner design](#). [4] Finly, [public automated-test run](#). [5] Finly, [live application](#), [public repository](#), [mathematical technical note](#), [engineering appendix](#), [demo video](#), and [pitch deck](#).